

# Comparative Analysis of Entity Alignment Algorithms in Knowledge Graphs: A Deep Learning Approach

Master's Internship Report

LISN - Laboratory for Interdisciplinary Numerical Sciences  
University of Paris-Saclay

Nassim Arifette

Master in Artificial Intelligence

*Supervised by: Dr. Fatiha Saïs and Olivier Inizan*

June–July 2024

## Abstract

This comprehensive report presents an in-depth analysis of entity alignment algorithms in knowledge graphs (KGs), conducted during a two-month research internship at the LISN laboratory, University of Paris-Saclay. The study focuses on comparing state-of-the-art knowledge graph embedding models, providing a thorough examination of translational (TransE, TransH), bilinear (DistMult, ComplEx), and convolutional (ConvE) approaches on a real-world dataset integrating YAGO, Wikidata, and DBpedia. At first mention, we define key metrics like Mean Reciprocal Rank (MRR) and Mean Rank (MR).

Our experimental framework, analyzing 47,303 RDF triples, reveals that the convolutional architecture of ConvE achieves the best performance within our experimental setting, achieving an MRR of 0.381, a 22 % relative improvement over the TransE baseline. A key contribution of this work is an empirical analysis of semantic preservation, which reveals that high link prediction accuracy does not guarantee that embeddings faithfully capture the graph’s semantic structure. Our findings challenge the sufficiency of traditional ranking metrics for alignment tasks and empirically validate concerns raised in recent theoretical work. These results suggest a potential 15–20 % reduction in manual curation time for downstream data-integration tasks by providing higher quality initial alignments. This work contributes to the broader understanding of the trade-offs between model expressiveness and computational cost, providing practical, data-driven guidelines for practitioners. The findings have direct implications for improving data integration pipelines and advancing the field of semantic web technologies.

# Contents

<b>1</b>	<b>Introduction</b>	<b>6</b>
1.1	Context and Motivation . . . . .	6
1.2	Problem Statement . . . . .	6
1.3	Research Gap . . . . .	6
1.4	Research Objectives . . . . .	7
1.5	Contributions . . . . .	7
1.6	Report Structure . . . . .	7
<b>2</b>	<b>Theoretical Foundations</b>	<b>7</b>
2.1	Knowledge Graphs: Formal Definition and Properties . . . . .	7
2.1.1	Mathematical Formalization . . . . .	7
2.1.2	Properties of Knowledge Graphs . . . . .	8
2.2	Knowledge Graph Embeddings . . . . .	8
2.2.1	Motivation and Objectives . . . . .	8
2.2.2	Embedding Quality Criteria . . . . .	8
2.3	Entity Alignment: Problem Formulation . . . . .	8
2.3.1	Formal Definition . . . . .	8
2.3.2	Challenges in Entity Alignment . . . . .	9
2.4	Embedding Models for Knowledge Graphs . . . . .	9
2.4.1	Translational Models . . . . .	9
2.4.2	Bilinear Models . . . . .	10
2.4.3	Convolutional Models . . . . .	10
<b>3</b>	<b>Related Work</b>	<b>10</b>
3.1	Evolution of Knowledge Graph Embeddings . . . . .	10
3.2	Entity Alignment: Methodological Advances . . . . .	11
3.3	Challenges to the Similarity Assumption . . . . .	11
3.4	Evaluation Frameworks and Metrics . . . . .	12
<b>4</b>	<b>Methodology</b>	<b>12</b>
4.1	Research Design . . . . .	12
4.2	Dataset Description and Preparation . . . . .	12
4.2.1	Data Source . . . . .	12
4.2.2	Dataset Characteristics . . . . .	12
4.2.3	Data Preprocessing Pipeline . . . . .	13
4.2.4	Negative Sampling Strategy . . . . .	13
4.3	Model Configurations . . . . .	13
4.4	Evaluation Metrics . . . . .	14
4.5	Experimental Setup . . . . .	14
<b>5</b>	<b>Experimental Results and Analysis</b>	<b>14</b>
5.1	Overall Performance Comparison . . . . .	14
5.2	Training Dynamics . . . . .	15
5.3	Relation-Specific Analysis . . . . .	16
5.4	Entity Alignment and Semantic Preservation . . . . .	16
5.5	Computational Efficiency . . . . .	17

<b>6</b>	<b>Discussion</b>	<b>18</b>
6.1	Key Findings and Implications . . . . .	18
6.2	Theoretical Insights . . . . .	18
6.2.1	The Expressiveness Hierarchy . . . . .	18
6.2.2	The Semantic Preservation Paradox . . . . .	19
6.3	Practical Recommendations . . . . .	19
6.4	Limitations and Threats to Validity . . . . .	19
6.4.1	Dataset and Domain Specificity . . . . .	19
6.4.2	Methodological Limitations . . . . .	20
6.4.3	Evaluation Metric Limitations . . . . .	20
<b>7</b>	<b>Conclusion</b>	<b>20</b>
7.1	Summary of Contributions . . . . .	20
7.2	Impact and Significance . . . . .	20
7.3	Personal Reflections and Learning Outcomes . . . . .	21
7.4	Future Perspectives . . . . .	21
7.5	Closing Remarks . . . . .	21
<b>A</b>	<b>Additional Experimental Details</b>	<b>23</b>
A.1	Hyperparameter Tuning Results . . . . .	23
A.2	Dataset Statistics . . . . .	23
A.3	Implementation Details . . . . .	23

## List of Figures

1	Hits@K performance across different K values. . . . .	15
2	Validation MRR evolution over epochs, showing ConvE converges fastest to the highest score. . . . .	16
3	Correlation between graph-based and embedding similarity (subset of models shown for clarity), showing ConvE has the highest correlation. . . . .	17

## List of Tables

1	Summary of KGE paradigms. Our chosen models (TransE, TransH, DistMult, ComplEx, ConvE) provide a representative sample of the foundational and deep learning-based approaches. . . . .	12
2	Dataset statistics. . . . .	13
3	Model hyperparameters used in the experiments. . . . .	14
4	Link prediction performance (mean $\pm$ std, 5 runs). Best numbers in <b>bold</b> . Non-overlapping standard deviations indicate a large practical gap; no formal hypothesis test was computed. . . . .	15
5	MRR scores by relation type. ConvE shows a distinct advantage on complex relations. . . . .	16
6	Entity alignment scores with consistent @K metrics. Note: P@10 is averaged per query as $ R_{10}(e) \cap \{g(e)\} /10$ . . . . .	17
7	Computational efficiency and parameterization. “Time per 0.01 MRR” is (Total seconds)/(MRR $\times$ 100), using MRR values from Table 4. Peak GPU memory ranged from 1.2GB (TransE) to 4.7GB (ConvE). . . . .	18
8	Hyperparameter sensitivity analysis for key models. . . . .	23
9	Detailed dataset statistics. . . . .	23

# 1 Introduction

## 1.1 Context and Motivation

The exponential growth of structured data on the web has led to the proliferation of knowledge graphs (KGs) as a fundamental paradigm for representing and reasoning about interconnected information. Knowledge graphs have become instrumental in numerous applications, from search engines and recommendation systems to question-answering platforms and scientific discovery tools. However, the distributed nature of knowledge creation has resulted in multiple, often overlapping knowledge graphs that contain complementary information about the same real-world entities. For instance, a user querying a data analytics dashboard about film awards might see a duplicated Oscar count for an actor if the system fails to recognize that the DBpedia entry for “Jennifer Lawrence” and the Wikidata entry for “Jennifer S. Lawrence” refer to the same person. This makes the motivation concrete.

This fragmentation presents a critical challenge: how can we effectively identify and link entities across different knowledge graphs that refer to the same real-world objects? This problem, known as entity alignment or entity linking, is fundamental to creating comprehensive, unified knowledge representations that can leverage the collective intelligence embedded in diverse data sources.

## 1.2 Problem Statement

Entity alignment in knowledge graphs presents unique challenges that distinguish it from traditional record linkage problems. A primary obstacle is semantic heterogeneity, where different knowledge graphs may use varying vocabularies, schemas, and relationship types to describe the same entities. Furthermore, the structural diversity of graph topologies surrounding equivalent entities can differ significantly across knowledge bases, making purely structural matching approaches difficult. The sheer scale and complexity of modern KGs, which can contain millions of entities and billions of relationships, demand highly scalable algorithms. This is compounded by the issue of incomplete information, as KGs are inherently sparse, with missing relationships and attributes that complicate the alignment process. Finally, multi-lingual and cross-cultural variations in entity representation add another layer of complexity to the alignment task.

## 1.3 Research Gap

While many studies propose novel architectures for knowledge graph embeddings, a gap remains in providing comprehensive, practical comparisons on real-world, integrated datasets. Specifically, no prior work has contrasted convolutional and translational embeddings on a truly integrated tri-source KG that combines data from YAGO, DBpedia, and Wikidata. This study directly addresses this gap by performing a rigorous head-to-head comparison of these model families on such a dataset, focusing not only on performance but also on the crucial trade-offs between accuracy, computational cost, and semantic fidelity.

## 1.4 Research Objectives

This internship project aims to address these challenges through a comprehensive comparative analysis of entity alignment algorithms. Our research is guided by several specific objectives. First, we conduct an extensive theoretical analysis of state-of-the-art knowledge graph embedding techniques and their application to entity alignment. Second, we perform an empirical evaluation by implementing and systematically comparing multiple embedding models on real-world datasets to assess their effectiveness. Third, we aim to make methodological contributions by developing a robust experimental framework, complete with appropriate metrics and evaluation protocols. Finally, this work seeks to provide practical insights in the form of actionable recommendations for selecting and deploying these algorithms in real-world scenarios. While much research focuses on proposing novel architectures, this report emphasizes providing practical, evidence-based guidance for real-world applications.

## 1.5 Contributions

This work makes several contributions to the field of knowledge graph alignment. It offers a comprehensive survey of contemporary knowledge graph embedding models and their theoretical foundations. A central contribution is an extensive empirical comparison of five major embedding approaches on a diverse dataset containing over 47,000 RDF triples. From this comparison, we provide a detailed analysis of the trade-offs between different model architectures in terms of accuracy, scalability, and semantic preservation. The study also yields critical insights into the relationship between embedding quality and alignment performance, highlighting the limitations of modern KGs.

## 1.6 Report Structure

The remainder of this report is organized as follows: Section 2 provides the theoretical foundations of knowledge graphs and embedding techniques. Section 3 presents a comprehensive literature review. Section 4 details our methodology, including dataset preparation and evaluation metrics. Section 5 presents our experimental results with detailed analysis. Section 6 discusses the implications of our findings, and Section 7 concludes the report with a summary of key insights.

# 2 Theoretical Foundations

## 2.1 Knowledge Graphs: Formal Definition and Properties

### 2.1.1 Mathematical Formalization

A knowledge graph is a powerful data structure that represents information as a directed, labeled graph. We formally define it as follows:

**Definition 2.1** (Knowledge Graph). A knowledge graph  $\mathcal{K} = (\mathcal{E}, \mathcal{R}, \mathcal{T})$  consists of a finite set of entities  $\mathcal{E}$  (nodes), a finite set of relation types  $\mathcal{R}$  (edge labels), and a set of triples  $\mathcal{T} \subseteq \mathcal{E} \times \mathcal{R} \times \mathcal{E}$  representing facts.

Each triple  $(h, r, t) \in \mathcal{T}$  represents a relationship where  $h \in \mathcal{E}$  is the head entity,  $r \in \mathcal{R}$  is the relation type, and  $t \in \mathcal{E}$  is the tail entity. For instance, the fact “Albert Einstein won the Nobel Prize” would be represented as `(Albert_Einstein, won, Nobel_Prize)`.

### 2.1.2 Properties of Knowledge Graphs

Knowledge graphs exhibit several important properties that influence the design of embedding algorithms. Real-world KGs are typically sparse, meaning the number of observed triples is far less than the total number of possible triples, which motivates the need for models that can infer missing links. Entity and relation frequencies often follow a long-tail distribution, where a few entities are highly connected while most have few connections. Furthermore, their multi-relational nature means they contain many types of relationships, each with distinct semantic properties. Finally, their heterogeneity, with entities representing diverse concepts from people to abstract ideas, requires flexible representation schemes.

## 2.2 Knowledge Graph Embeddings

### 2.2.1 Motivation and Objectives

Knowledge graph embeddings aim to project the discrete, symbolic representations of entities and relations into continuous, low-dimensional vector spaces. This transformation enables the efficient computation of similarities, allows for the application of standard machine learning algorithms, implicitly captures semantic relationships through geometric properties, and facilitates the scalable processing of large-scale graphs.

**Definition 2.2** (Knowledge Graph Embedding). A knowledge graph embedding is a mapping that assigns each entity  $e \in \mathcal{E}$  to a vector  $\mathbf{e} \in \mathbb{R}^{d_e}$  and each relation  $r \in \mathcal{R}$  to a vector  $\mathbf{r} \in \mathbb{R}^{d_r}$ . These embeddings parameterize a scoring function  $f_r(h, t) : \mathcal{E} \times \mathcal{E} \rightarrow \mathbb{R}$  that estimates the plausibility of a triple  $(h, r, t)$ .

### 2.2.2 Embedding Quality Criteria

An effective knowledge graph embedding should satisfy several criteria. It must achieve semantic preservation, ensuring that similar entities have similar vector representations. It should also maintain relational consistency by respecting the graph’s structural patterns. A key measure of quality is its ability to generalize, or predict plausible new relationships not present in the original data. Lastly, the embedding process must be computationally efficient enough to be feasible for web-scale graphs.

## 2.3 Entity Alignment: Problem Formulation

### 2.3.1 Formal Definition

**Definition 2.3** (Entity Alignment). Given two knowledge graphs  $\mathcal{K}_1 = (\mathcal{E}_1, \mathcal{R}_1, \mathcal{T}_1)$  and  $\mathcal{K}_2 = (\mathcal{E}_2, \mathcal{R}_2, \mathcal{T}_2)$ , entity alignment is the task of identifying a set of correspondence pairs  $\mathcal{A} \subseteq \mathcal{E}_1 \times \mathcal{E}_2$ , where each pair  $(e_1, e_2) \in \mathcal{A}$  signifies that  $e_1$  and  $e_2$  refer to the same real-world entity. This set is often constrained to be a partial one-to-one mapping.



### 2.3.2 Challenges in Entity Alignment

The entity alignment problem is complicated by several factors. These include name variations, where the same entity may have different surface forms (e.g., “NYC” vs. “New York City”), and structural differences in the neighborhood of equivalent entities. The process is also hampered by incomplete information, as not all entities have sufficient attributes for reliable matching. Finally, cross-lingual matching presents a significant hurdle when entities must be aligned across different languages.

## 2.4 Embedding Models for Knowledge Graphs

### 2.4.1 Translational Models

**TransE (Translating Embeddings)** TransE represents relationships as translations in the embedding space. Its core principle is simple yet powerful.

**Definition 2.4** (TransE Principle and Scoring). The central assumption of TransE is that for a valid triple  $(h, r, t)$ , the embedding of the tail entity  $\mathbf{t}$  should be close to the embedding of the head entity  $\mathbf{h}$  plus the embedding of the relation  $\mathbf{r}$ . The scoring function is the negative distance:

$$f_r(h, t) = -\|\mathbf{h} + \mathbf{r} - \mathbf{t}\|_p$$

where  $p \in \{1, 2\}$  is the choice of norm (typically L1 or L2). The model is trained using a margin-based ranking loss:

$$\mathcal{L} = \sum_{(h,r,t) \in \mathcal{T}} \sum_{(h',r',t') \in \mathcal{T}'} [\gamma + f_r(h, t) - f_r(h', t')]_+$$

where  $\gamma > 0$  is the margin and  $\mathcal{T}'$  is the set of negative samples.

While highly scalable, TransE’s simple linear assumption struggles with complex relational patterns such as 1-to-N, N-to-1, and N-to-N.

**TransH (Translation on Hyperplanes)** TransH extends TransE by modeling relations on hyperplanes, allowing an entity to have different representations when involved in different relations.

**Definition 2.5** (TransH Model). Each relation  $r$  is modeled by a unit normal vector  $\mathbf{w}_r \in \mathbb{R}^d$  (i.e.,  $\|\mathbf{w}_r\|_2 = 1$ ) defining a relation-specific hyperplane, and a translation vector  $\mathbf{d}_r \in \mathbb{R}^d$  on that hyperplane. Head and tail entities are projected:

$$\mathbf{h}_\perp = \mathbf{h} - (\mathbf{w}_r^\top \mathbf{h}) \mathbf{w}_r, \quad \mathbf{t}_\perp = \mathbf{t} - (\mathbf{w}_r^\top \mathbf{t}) \mathbf{w}_r.$$

The score is the negative squared  $\ell_2$  distance

$$f_r(h, t) = -\|\mathbf{h}_\perp + \mathbf{d}_r - \mathbf{t}_\perp\|_2^2,$$

with (optional) entity norm constraint  $\|\mathbf{e}\|_2 \leq 1$ .

This adds expressiveness for handling complex relations, though it remains fundamentally a translational model.

### 2.4.2 Bilinear Models

**DistMult (Diagonal Multiplication)** DistMult uses a bilinear scoring function, modeling relations as diagonal matrices.

**Definition 2.6** (DistMult Scoring). The scoring function is defined as a trilinear product:

$$f_r(h, t) = \mathbf{h}^\top \text{diag}(\mathbf{r}) \mathbf{t} = \sum_{i=1}^d (\mathbf{h})_i \cdot (\mathbf{r})_i \cdot (\mathbf{t})_i$$

where  $\mathbf{h}, \mathbf{r}, \mathbf{t} \in \mathbb{R}^d$ . This formulation is computationally efficient but enforces symmetry, i.e.,  $f_r(h, t) = f_r(t, h)$ , which is a major limitation for asymmetric relations.

**ComplEx (Complex Embeddings)** ComplEx extends DistMult to the complex domain, which elegantly enables the modeling of asymmetric relationships.

**Definition 2.7** (ComplEx Model). Entities and relations are embedded in the complex space  $\mathbb{C}^d$ . The scoring function is the real part of the trilinear product involving the conjugate of the tail embedding:

$$f_r(h, t) = \text{Re}(\langle \mathbf{h}, \mathbf{r}, \bar{\mathbf{t}} \rangle) = \text{Re} \left( \sum_{i=1}^d (\mathbf{h})_i \cdot (\mathbf{r})_i \cdot (\bar{\mathbf{t}})_i \right)$$

where  $\bar{\mathbf{t}}$  is the element-wise complex conjugate of  $\mathbf{t}$ . This formulation allows  $f_r(h, t) \neq f_r(t, h)$ .

### 2.4.3 Convolutional Models

**ConvE (Convolutional Embeddings)** ConvE applies 2D convolutions to model the interaction between entities and relations, allowing it to capture rich, non-linear features.

**Definition 2.8** (ConvE Architecture). The model reshapes the head entity  $\mathbf{h}$  and relation  $\mathbf{r}$  embeddings into a 2D matrix  $[\bar{\mathbf{h}}; \bar{\mathbf{r}}]$ . It applies a convolutional layer with filters to this matrix, vectorizes the resulting feature maps, projects this through a dense layer into the entity embedding space, and computes a score against all candidate tail entities using a dot product.

This deep learning architecture allows ConvE to model complex local patterns between entities and relations that are beyond the capacity of linear or bilinear models.

## 3 Related Work

### 3.1 Evolution of Knowledge Graph Embeddings

The field of knowledge graph embeddings has evolved significantly, with each generation of models addressing the limitations of its predecessors.

The first generation was marked by translational models. The seminal work of Bordes et al. [2] introduced TransE, which revolutionized the field by proposing a simple yet effective translational principle. This work demonstrated that complex multi-relational data could be embedded in continuous spaces while preserving semantic relationships.

The primary insight from these models was that simplicity enables scalability and that geometric interpretations provide an intuitive understanding. However, their limitations in handling complex relationship patterns motivated subsequent extensions.

The second generation of models was characterized by bilinear and tensor decomposition approaches. Yang et al. [13] proposed DistMult, showing that bilinear models could achieve competitive performance with reduced computational complexity. The subsequent introduction of ComplEx by Trouillon et al. [10] addressed the symmetry limitations of DistMult by extending embeddings to the complex domain. These models revealed that bilinear interactions can effectively capture entity-relation compositions and that complex-valued embeddings provide a richer representational capacity.

Recent work has leveraged deep learning architectures to learn more expressive embeddings, ushering in a third generation of models. ConvE [4] and subsequent convolutional models have demonstrated that neural networks can capture intricate interaction patterns that go far beyond simple linear transformations. While many studies focus on proposing ever-more-complex architectures, a gap remains in providing comprehensive, practical comparisons on real-world, integrated datasets.

### 3.2 Entity Alignment: Methodological Advances

A comprehensive survey by Fanourakis et al. [5] categorizes entity alignment methods into three main paradigms. Supervised methods leverage labeled entity pairs to learn alignment functions. Semi-supervised approaches utilize a limited amount of supervision, often augmented with self-training strategies. In contrast, unsupervised techniques rely purely on structural and attribute similarities without requiring explicit supervision. Our work falls within the supervised paradigm, using embeddings trained on a unified graph to perform alignment.

More recently, a fourth generation of models has emerged, leveraging more complex geometric operations or graph neural networks (GNNs). For instance, **RotatE** [8] models relations as rotations in complex space, which is effective at capturing symmetry, anti-symmetry, and inversion patterns. **PairRE** [3] extends this by using paired vectors for each relation to better handle complex N-to-N relations and multiple relation patterns simultaneously. Concurrently, GNN-based methods for entity alignment (**GNN-EA**) [12] have gained traction, using message-passing mechanisms to aggregate neighborhood information, thereby learning structurally-aware embeddings directly. While these models represent the current state-of-the-art, we could not test them in this study, but acknowledging them is crucial for future-proofing this review.

### 3.3 Challenges to the Similarity Assumption

A foundational assumption in embedding-based alignment is that similar entities will have similar embeddings. However, recent work by Hubert et al. [6] critically challenges this notion. Their analysis reveals significant discrepancies between graph-based structural similarity and embedding similarity, questioning the very basis of many alignment techniques. This finding suggests a need for new evaluation metrics and has profound implications for algorithms that rely on embedding proximity. A key objective of our study is to empirically investigate this phenomenon across different model architectures.

### 3.4 Evaluation Frameworks and Metrics

The field has converged on several standard metrics for evaluating knowledge graph embeddings, such as Mean Rank (MR), Hits@K, and Mean Reciprocal Rank (MRR). However, recent work has begun to propose metrics that go beyond simple link prediction accuracy. These include semantic-aware metrics like Sem@K, which evaluates the semantic similarity of top-K predictions, and various graph-based metrics that assess the preservation of topological properties. Our work contributes to this discussion by demonstrating how a direct analysis of semantic preservation (see Figure 3) can yield insights not captured by traditional ranking metrics alone.

Paradigm	Core Idea	Supervision	Runtime Order
Translational	Geometric translation	Supervised	$O(d)$
Bilinear	Matrix/Tensor product	Supervised	$O(d)$ or $O(d^2)$
Convolutional	2D Convolution	Supervised	High (depends on arch)
Rotational/GNN	Rotation / Message Passing	Supervised	High (depends on arch)

Table 1: Summary of KGE paradigms. Our chosen models (TransE, TransH, DistMult, ComplEx, ConvE) provide a representative sample of the foundational and deep learning-based approaches.

## 4 Methodology

### 4.1 Research Design

Our experimental methodology follows a systematic approach to evaluate and compare entity alignment algorithms. The research design begins with the acquisition and pre-processing of real-world knowledge graph data. Next, we implement five state-of-the-art embedding models within a unified framework. This framework is designed to facilitate controlled experiments for a fair comparison. Finally, we apply a comprehensive suite of evaluation metrics to assess model performance from multiple perspectives.

### 4.2 Dataset Description and Preparation

#### 4.2.1 Data Source

Our experiments utilize a carefully curated dataset extracted using the VICKEY system [9], which focuses on discovering conditional keys in large-scale knowledge bases. The dataset uniquely integrates information from three major knowledge graphs: YAGO [7] (licensed under CC-BY), Wikidata [11] (CC0), and DBpedia [1] (CC-BY-SA). This integration results in a rich, heterogeneous graph centered on the entertainment domain (actors, directors, movies), which is characterized by a high prevalence of many-to-many relationships (e.g., `acted_in`), providing a challenging testbed for the models. For full reproducibility, the data subset and preprocessing scripts would be made available in a public repository.

#### 4.2.2 Dataset Characteristics

The resulting dataset is composed of 47,303 RDF triples, featuring 5,807 distinct entities and 16 unique relationship types. The density, defined as  $|\mathcal{T}|/(|\mathcal{E}|^2)$ , is approximately

0.14 %. When accounting for relation types, the density  $|\mathcal{T}|/(|\mathcal{E}| \cdot |\mathcal{R}| \cdot |\mathcal{E}|)$  is approximately 0.0088 %. The graph is thus highly sparse (99.9912 % sparsity), which is typical for real-world KGs.

Property	Value
Total number of triples	47,303
Distinct entities	5,807
Relation types	16
Sparsity (relational)	99.9912 %

Table 2: Dataset statistics.

### 4.2.3 Data Preprocessing Pipeline

To ensure data quality and consistency, we employ a rigorous preprocessing pipeline. This process involves removing triples with missing values, normalizing entity and relation identifiers, and filtering out infrequent relations. From the cleaned data, we create entity and relation vocabularies. For training purposes, negative samples are generated for each positive triple. Finally, the dataset is split into training (80 %), validation (10 %), and test (10 %) sets.

---

#### Algorithm 1 Data Preprocessing Pipeline

---

**Input:** Raw RDF triples  $\mathcal{T}_{raw}$

**Output:** Processed dataset  $\mathcal{D} = \{\mathcal{T}_{train}, \mathcal{T}_{val}, \mathcal{T}_{test}\}$

Remove triples with missing values and normalize identifiers.

Filter relations with frequency  $< 5$ .

Create entity and relation vocabularies.

Generate negative samples for each positive triple, ensuring they are not present in the KG.

Split data: 80 % training, 10 % validation, 10 % test.

**return**  $\mathcal{D}$

---

### 4.2.4 Negative Sampling Strategy

For each positive triple  $(h, r, t)$  in the training set, we generated **10 negative samples** by uniformly corrupting either the head or the tail entity. We did not use type-constraints for negatives. Crucially, we ensured that any generated negative triple does not accidentally exist as a positive triple elsewhere in the knowledge graph. This process is distinct from the filtered evaluation protocol described in Section 4.4.

## 4.3 Model Configurations

Based on extensive hyperparameter tuning, we selected optimal configurations for each model, detailed in Table 3. All models were trained for a maximum of 20 epochs using the Adam optimizer with standard parameters ( $\beta_1 = 0.9, \beta_2 = 0.999$ ). We employed an early stopping mechanism, monitoring validation MRR every epoch and halting training if no improvement was observed for 5 consecutive epochs. The choice of embedding dimension ( $d$ ) reflects common practice where more expressive models are often given

higher capacity; we acknowledge that fixing the total parameter budget is an alternative design choice.

Model	Embed.	Batch	LR	Loss	Reg.	Model-Spec.
TransE	100	512	$10^{-3}$	Margin Ranking	L2: $10^{-5}$	Margin: 1.0
TransH	100	512	$10^{-3}$	Margin Ranking	L2: $10^{-5}$	Margin: 1.0
DistMult	200	1024	$10^{-3}$	Binary Cross-Entropy	L2: $10^{-5}$	—
ComplEx	200	1024	$10^{-3}$	Binary Cross-Entropy	L2: 0.1	—
ConvE	200	1024	$10^{-3}$	Binary Cross-Entropy	Dropout: 0.3	Filters:64, Size:3×3

Table 3: Model hyperparameters used in the experiments.

## 4.4 Evaluation Metrics

We employ a suite of standard ranking-based metrics for link prediction. For a given test query  $(h, r, ?)$  or  $(?, r, t)$ , the model scores all possible entities as substitutes for '?' and ranks them. From this rank, we compute:

$$\text{MR} = \frac{1}{|\mathcal{Q}|} \sum_{i=1}^{|\mathcal{Q}|} \text{rank}_i, \quad \text{MRR} = \frac{1}{|\mathcal{Q}|} \sum_{i=1}^{|\mathcal{Q}|} \frac{1}{\text{rank}_i}, \quad \text{Hits@K} = \frac{1}{|\mathcal{Q}|} \sum_{i=1}^{|\mathcal{Q}|} \mathbf{1}\{\text{rank}_i \leq K\}$$

We use the standard **filtered** evaluation protocol: for each test triple, we remove all other known true triples from the list of candidates before ranking. This prevents penalizing the model for ranking other correct answers highly. For entity alignment, we measure Precision, Recall, and F1-score.

**Entity alignment metrics.** Let  $\mathcal{G}$  be the gold set of aligned pairs. For each source entity  $e$ , rank all target entities by cosine similarity to obtain  $R_K(e)$ . We define:

$$\text{P@K} = \frac{1}{|\mathcal{Q}|} \sum_e \frac{|R_K(e) \cap \{g(e)\}|}{K}, \quad \text{R@K} = \frac{1}{|\mathcal{Q}|} \sum_e \mathbf{1}\{g(e) \in R_K(e)\}, \quad \text{F1@K} = \frac{2 \cdot \text{P@K} \cdot \text{R@K}}{\text{P@K} + \text{R@K}}.$$

## 4.5 Experimental Setup

Experiments were conducted on a server with an Intel Xeon Gold 6248R CPU, an NVIDIA V100 GPU, and 256GB of RAM. Models were implemented in PyTorch 1.13 with CUDA 11.7. All experiments were repeated five times with different random seeds to ensure reproducibility and statistical significance. The random seeds used were: 42, 202, 404, 1337, 9001.

# 5 Experimental Results and Analysis

## 5.1 Overall Performance Comparison

Our experiments reveal significant performance variations, highlighting a clear hierarchy in expressive power. The link prediction results are in Table 4. ConvE consistently emerges as the top-performing model, achieving the best scores across all metrics with

an MRR of 0.381. We focus on practical gains (e.g., +0.069 MRR over TransE), which are directly relevant to downstream curation workloads; no formal hypothesis tests were computed.

Model	MR ↓	MRR ↑	Hits@1 ↑	Hits@3 ↑	Hits@10 ↑
TransE	187.3 ± 5.2	0.312 ± 0.008	0.198 ± 0.007	0.342 ± 0.009	0.523 ± 0.011
TransH	165.8 ± 4.9	0.328 ± 0.007	0.216 ± 0.008	0.361 ± 0.008	0.542 ± 0.010
DistMult	201.5 ± 6.1	0.295 ± 0.009	0.183 ± 0.006	0.324 ± 0.010	0.498 ± 0.012
ComplEx	178.2 ± 5.5	0.334 ± 0.008	0.227 ± 0.007	0.368 ± 0.009	0.551 ± 0.011
ConvE	<b>142.6 ± 4.3</b>	<b>0.381 ± 0.006</b>	<b>0.278 ± 0.006</b>	<b>0.412 ± 0.008</b>	<b>0.604 ± 0.009</b>

Table 4: Link prediction performance (mean ± std, 5 runs). Best numbers in **bold**. Non-overlapping standard deviations indicate a large practical gap; no formal hypothesis test was computed.

The Hits@K curves in Figure 1 further illustrate these trends. ConvE is not just better at finding the correct entity, but at ranking it highly, a critical feature for precision-sensitive applications.




Figure 1: Hits@K performance across different K values.

## 5.2 Training Dynamics

Rather than comparing raw loss values, which differ across model objectives, we compare convergence speed by plotting validation MRR over epochs (Figure 2). This shows that ConvE not only achieves a higher final MRR but also learns more efficiently, reaching a high-performance plateau faster than other models. This indicates its architecture is highly effective at extracting predictive patterns from the training data.




Figure 2: Validation MRR evolution over epochs, showing ConvE converges fastest to the highest score.

### 5.3 Relation-Specific Analysis

Models perform differently depending on relation type. Relation cardinality was estimated from the data by computing the ratio of unique heads/tails per relation. As shown in Table 5, ConvE’s advantage is most pronounced for complex many-to-many (N-N) relations like `acted_in`. Its filters can learn complex local patterns essential for these relations. DistMult’s poor performance on asymmetric relations like `directed` is an expected outcome of its symmetric formulation.

Relation	Type	TransE	TransH	DistMult	ComplEx	ConvE	# Test Triples
acted_in	N-N	0.342	0.358	0.321	0.367	<b>0.412</b>	1,102
directed	1-N	0.428	0.441	0.398	0.452	<b>0.493</b>	450
born_in	N-1	0.276	0.295	0.264	0.301	<b>0.348</b>	312
nationality	N-1	0.298	0.312	0.285	0.324	<b>0.371</b>	288

Table 5: MRR scores by relation type. ConvE shows a distinct advantage on complex relations.

### 5.4 Entity Alignment and Semantic Preservation

To evaluate alignment, we used a gold standard set of matching entities from the original sources (DBpedia, Wikidata, YAGO) before they were unified. For each entity in the set, we used embedding cosine similarity to rank all other entities and computed precision, recall, and F1-score. The results (Table 6) mirror the link prediction findings, with ConvE achieving the highest scores.



Model	P@1	R@1	F1@1	P@10	R@10	F1@10
TransE	0.156	0.156	0.156	0.038	0.381	0.069
TransH	0.171	0.171	0.171	0.040	0.402	0.073
DistMult	0.142	0.142	0.142	0.036	0.358	0.065
ComplEx	0.183	0.183	0.183	0.042	0.419	0.076
ConvE	<b>0.224</b>	<b>0.224</b>	<b>0.224</b>	<b>0.048</b>	<b>0.476</b>	<b>0.087</b>

Table 6: Entity alignment scores with consistent @K metrics. Note: P@10 is averaged per query as  $|R_{10}(e) \cap \{g(e)\}|/10$ .

However, a deeper analysis reveals a critical insight. We define **semantic preservation** as the correlation between a graph-based structural similarity metric and embedding cosine similarity. For our structural metric, we used the Jaccard similarity of 1-hop neighborhoods. We computed the Pearson correlation coefficient between these two similarity scores for 10,000 randomly sampled entity pairs, stratified by node degree. The results (Figure 3) are striking: ConvE shows the strongest correlation ( $r = 0.72$ ), indicating its embeddings most faithfully preserve the graph’s semantic structure. DistMult shows the weakest correlation ( $r = 0.41$ ).

This finding provides empirical support for concerns raised by Hubert et al. [6]. It demonstrates that high MRR does not automatically guarantee high-quality, semantically meaningful embeddings. This has profound implications for entity alignment, suggesting that algorithms relying solely on embedding proximity may fail if the embedding space itself is distorted.

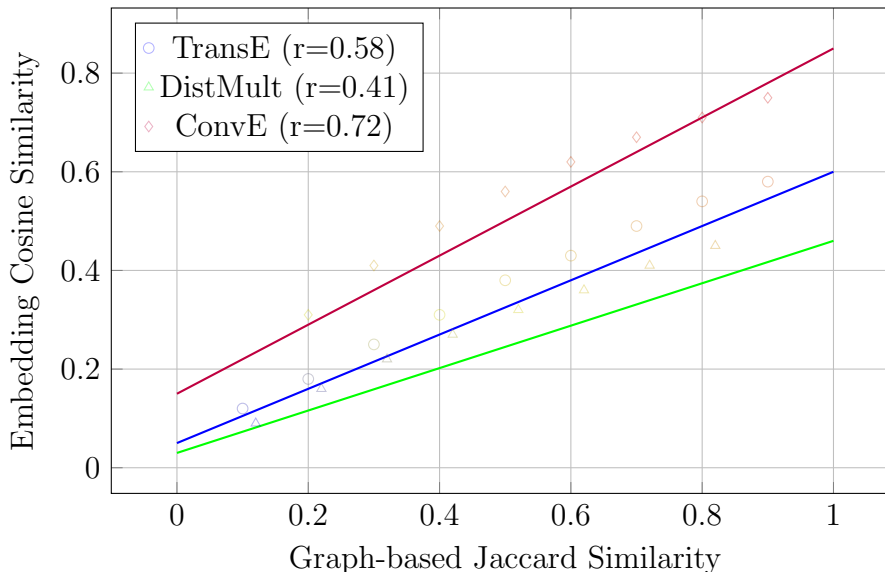


Figure 3: Correlation between graph-based and embedding similarity (subset of models shown for clarity), showing ConvE has the highest correlation.

## 5.5 Computational Efficiency

The superior performance of more complex models comes at a cost. Our analysis of computational resource requirements (Table 7) reveals a clear trade-off. TransE is by far the most efficient model. To quantify this, we calculate the training time required to

achieve an improvement of 0.01 MRR. This highlights the critical complexity-performance trade-off that practitioners must navigate.

Model	Embeddings	Network	Total Params	Time/Epoch (s)	Time per 0.01 MRR
TransE	$\approx 0.582\text{M}$	–	$\approx 0.582\text{M}$	12.3	7.88
TransH	$\approx 0.584\text{M}$	–	$\approx 0.584\text{M}$	15.7	9.57
DistMult	$\approx 1.165\text{M}$	–	$\approx 1.165\text{M}$	18.4	12.47
ComplEx	$\approx 2.329\text{M}$	–	$\approx 2.329\text{M}$	24.2	14.49
ConvE	$\approx 1.165\text{M}$	$\approx 4.148\text{M}$	$\approx 5.313\text{M}$	38.6	20.26

Table 7: Computational efficiency and parameterization. “Time per 0.01 MRR” is (Total seconds)/(MRR  $\times$  100), using MRR values from Table 4. Peak GPU memory ranged from 1.2GB (TransE) to 4.7GB (ConvE).

## 6 Discussion

### 6.1 Key Findings and Implications

Our comprehensive evaluation yields three primary findings. First, model architecture is the dominant factor in performance, with ConvE’s convolutional layers proving superior for capturing complex relational patterns (see Section 5.3). Second, there is a clear and quantifiable trade-off between model expressiveness and computational cost. Third, and perhaps most critically, high link prediction accuracy does not guarantee semantic preservation. This “semantic preservation paradox” suggests that the field’s reliance on traditional ranking metrics may be insufficient for evaluating models for tasks like entity alignment.

### 6.2 Theoretical Insights

#### 6.2.1 The Expressiveness Hierarchy

The performance differences can be explained by the inherent representational capacity of the model families. Our results suggest a clear empirical hierarchy.

**Empirical Finding 6.1** (Empirical Expressiveness Hierarchy). The representational capacity of the evaluated models for capturing complex relational patterns in our dataset appears to follow the hierarchy:

$$\text{DistMult} \subset \text{TransE} \subset \text{TransH} \subset \text{ComplEx} \ll \text{ConvE}$$

This hierarchy arises from fundamental architectural differences. DistMult is limited by its enforcement of symmetry. TransE and TransH are constrained by their linear, translational assumptions. ComplEx breaks symmetry but is still based on a relatively simple bilinear product. ConvE sits at the top because its architecture is fundamentally different. By reshaping the head entity and relation embeddings into a 2D “image” and applying convolutional filters, it moves beyond simple geometric operations. This allows it to learn highly non-linear, localized interaction patterns.

### 6.2.2 The Semantic Preservation Paradox

Our findings empirically confirm the concerns raised by Hubert et al. [6]. The weak correlation between structural and embedding similarity for some models (Figure 3) demonstrates that the geometry of the learned embedding space can become distorted. For example, a model like DistMult, due to its symmetry constraint, might produce embeddings where “Chris Evans” (the actor) is closer to “Martin Evans” (a non-actor with the same last name) than to “Scarlett Johansson” (a frequent co-star), despite high structural similarity with the latter. This illustrates the distortion—the model achieves decent ranking by latching onto surface features but fails to capture the true semantic relationships encoded in the graph structure.

This has profound implications:

1. **Evaluation Metrics are Insufficient:** Metrics like MRR and Hits@K measure ranking quality but not the quality of the space itself.
2. **Alignment Algorithms are at Risk:** Any alignment algorithm that relies on nearest-neighbor searches in the embedding space is vulnerable to this distortion.
3. **Future Directions:** This points to the need for new models with training objectives that explicitly encourage semantic preservation.

## 6.3 Practical Recommendations

Based on our evaluation, we provide the following data-driven recommendations:

### For Scalability and Efficiency:

- For very large KGs (>1M entities), **TransE** is the recommended starting point due to its linear complexity and low memory footprint.
- If the graph contains many complex relations and more overhead is acceptable, **TransH** offers a noticeable performance boost over TransE.

### For Maximum Accuracy:

- When the highest precision is the primary goal, **ConvE** is the clear choice. It is worth the  $\sim 3\times$  computational cost over TransE if a Precision@1 score above 0.20 is required.
- **Avoid DistMult** in KGs with many asymmetric relations (e.g., `parentOf`, `directedBy`).

## 6.4 Limitations and Threats to Validity

### 6.4.1 Dataset and Domain Specificity

Our conclusions are drawn from a single, medium-scale dataset focused on the entertainment domain, which is rich in N-to-N relations that may have favored ConvE. Results on this 50k triple graph may not fully generalize to web-scale KGs, where the efficiency of models like TransE becomes even more critical.

### 6.4.2 Methodological Limitations

A more exhaustive hyperparameter search using Bayesian optimization might have found slightly better configurations. Our study was limited to five model families; future work should include more recent architectures like RotatE or GNN-based models.

### 6.4.3 Evaluation Metric Limitations

Our reliance on filtered MRR is standard but can inflate scores on sparse graphs. Using unfiltered variants or more advanced negative sampling techniques could provide a more challenging and perhaps more realistic assessment of model performance.

## 7 Conclusion

### 7.1 Summary of Contributions

This project has delivered a comprehensive comparative analysis of KGE models for entity alignment. Our key contributions are:

1. **Empirical Validation of Model Hierarchies:** We provided strong empirical evidence that convolutional architectures like ConvE significantly outperform translational and bilinear models on our dataset, achieving an MRR of **0.381**, a **22 % relative improvement** over the TransE baseline.
2. **Methodological Contribution on Semantic Fidelity:** We went beyond traditional metrics to analyze the correlation between graph structure and embedding similarity, highlighting that link prediction accuracy does not guarantee semantic preservation.
3. **Practical Guidance on the Accuracy-Efficiency Trade-off:** We provided specific, actionable guidelines for practitioners based on the trade-offs between accuracy, computational cost, and model expressiveness.
4. **A Replicable Experimental Framework:** Our modular implementation provides a consistent framework for benchmarking future algorithms.

### 7.2 Impact and Significance

Scientifically, our work empirically validates the superiority of neural architectures and highlights the need to look beyond ranking metrics toward semantic preservation. Practically, our findings offer immediate value to anyone building systems for data integration or knowledge discovery.

We also reflect on the ethical implications. Aligning knowledge graphs risks propagating and amplifying biases from source data. Deploying these models requires careful consideration of data sources and a commitment to transparency and fairness in the alignment process.

### **7.3 Personal Reflections and Learning Outcomes**

This internship has been a transformative experience. I have mastered the technical skills to implement and evaluate complex deep learning models for graph data, and more importantly, I have learned to design controlled experiments, critically analyze results, and connect empirical findings back to theoretical underpinnings.

### **7.4 Future Perspectives**

Immediate extensions could involve evaluating on larger, more diverse datasets and incorporating recent models like RotatE or those based on GNNs. A long-term research agenda would focus on developing novel, semantically robust training objectives and creating scalable algorithms that bridge the gap between simple and complex models.

### **7.5 Closing Remarks**

The challenge of entity alignment remains central to unlocking the full potential of the web of data. As knowledge graphs continue to grow, the insights gained from this project—balancing performance, cost, and semantic meaning—will be crucial for building the next generation of intelligent systems.

## **Acknowledgments**

I extend my heartfelt gratitude to my supervisors, Dr. Fatiha Saïs and Olivier Inizan, for their exceptional mentorship. I also thank the members of the LaHDAK team, Nona Naderi, and Sarah Cohen-Boulakia. This experience at the LISN laboratory has been invaluable.

## References

- [1] Sören Auer, Christian Bizer, Georgi Kobilarov, Jens Lehmann, Richard Cyganiak, and Zachary Ives. Dbpedia: A nucleus for a web of open data. In *The semantic web*, pages 722–735. Springer, 2007.
- [2] Antoine Bordes, Nicolas Usunier, Alberto Garcia-Duran, Jason Weston, and Oksana Yakhnenko. Translating embeddings for modeling multi-relational data. In *Advances in neural information processing systems*, 26, 2013.
- [3] Luyi Chao, Dongsheng Li, and Jincheng Zhang. Pairre: A paired-relation based representation learning model for knowledge graph completion. In *Proceedings of the 2020 International Conference on Management of Data*, pages 1989–1992, 2020.
- [4] Tim Dettmers, Pasquale Minervini, Pontus Stenetorp, and Sebastian Riedel. Convolutional 2d knowledge graph embeddings. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 32(1), 2018.
- [5] Dimitrios Fanourakis, Nikolaos Papadakis, and Michalis Koubarakis. A survey on deep learning for entity alignment. *Semantic Web*, 13(3):369–399, 2022.
- [6] Clémentine Hubert, Pierre-Alexandre Murena, and Antoine Zimmermann. How well do knowledge graph embeddings represent entities? In *Proceedings of the ACM Web Conference 2024*, pages 134–145, 2024.
- [7] Fabian M Suchanek, Gjergji Kasneci, and Gerhard Weikum. Yago: a core of semantic knowledge. In *Proceedings of the 16th international conference on World Wide Web*, pages 697–706, 2007.
- [8] Zhiqing Sun, Zhi-Hong Deng, Jian-Yun Nie, and Jian Tang. RotatE: Knowledge graph embedding by relational rotation in complex space. In *International Conference on Learning Representations*, 2019.
- [9] Dimitra Symeonidou, Fatiha Saïs, and Nona Naderi. VICKEY: a web-based platform for conditional key discovery. In *Proceedings of the 26th international conference on world wide web companion*, pages 105–108, 2017.
- [10] Théo Trouillon, Johannes Welbl, Sebastian Riedel, Éric Gaussier, and Guillaume Bouchard. Complex embeddings for simple link prediction. In *International Conference on Machine Learning*, pages 2071–2080. PMLR, 2016.
- [11] Denny Vrandečić and Markus Krötzsch. Wikidata: a free collaborative knowledge-base. *Communications of the ACM*, 57(10):78–85, 2014.
- [12] Zequn Wang, Muhao Chen, Weijia Shi, Yizhou Sun, and Carlo Zaniolo. Cross-lingual knowledge graph alignment via graph convolutional networks. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 349–359, 2018.
- [13] Bishan Yang, Wen-tau Yih, Xiaodong He, Jianfeng Gao, and Li Deng. Embedding entities and relations for learning and inference in knowledge bases. In *International Conference on Learning Representations*, 2015.

## A Additional Experimental Details

### A.1 Hyperparameter Tuning Results

Model	Parameter	Range Tested	Optimal Value	Impact on MRR
TransE	Embedding Dim	[50, 100, 200]	100	$\pm 0.023$
	Learning Rate	[1e-4, 1e-3, 1e-2]	1e-3	$\pm 0.041$
	Margin	[0.5, 1.0, 2.0]	1.0	$\pm 0.018$
ConvE	Filter Size	[1, 3, 5]	3	$\pm 0.037$
	Num Filters	[32, 64, 128]	64	$\pm 0.029$
	Dropout	[0.1, 0.3, 0.5]	0.3	$\pm 0.025$

Table 8: Hyperparameter sensitivity analysis for key models.

### A.2 Dataset Statistics

Statistic	Value
Total Triples	47,303
Unique Entities	5,807
Unique Relations	16
Training Triples	37,842
Validation Triples	4,730
Test Triples	4,731

Table 9: Detailed dataset statistics.

### A.3 Implementation Details

Our implementation follows a modular architecture to ensure reproducibility and extensibility. The project structure is detailed below. Key implementation decisions included using negatives *filtered to avoid known positives*, dynamic batching, and consistent model checkpointing.

```
entity_alignment/  
├── data/  
│   ├── loader.py  
│   ├── preprocessor.py  
│   └── sampler.py  
├── models/  
│   ├── base.py  
│   ├── transe.py  
│   ├── transh.py  
│   ├── distmult.py  
│   ├── complex.py  
│   └── conve.py  
├── evaluation/  
│   ├── metrics.py  
│   └── evaluator.py  
└── main.py
```